# Types

Our approach to types is to create a language that is *strongly static typed*. We add syntax to provide type information for values, and every program is checked for proper type matching prior to evaluating the program.

A *static typed* language is one in which type information can be determined at "compile time" rather than at runtime. In this way, any type errors are found before the program is evaluated. Saying that a language is *strongly typed* means that there are no type "holes" in the system: the declared type of a symbol (a variable or procedure parameter) dictates the types of the values that can be bound to the symbol. Everything is type checked, without exception, and all expressions must conform to type rules.

Our typed language is based on the `SET` language. In the `SET` language, only the integer value `0` is considered false, and all other values – including `ProcVals` – are true when used in conditional expressions. The Java programming language, which is staticaly typed, requires that the test expression in an `if` statement evaluate to a boolean. We follow this example in our typed language by introducing *primitive types*, `int` and `bool`, corresponding to integer and boolean expressed values, respectively. To implement this, we add a new `BoolVal` subclass of the `Val` class. With this addition, we can require that the test expression in an `if` expression be a boolean.

# Types

Every `IntVal` has type `int`, and every `BoolVal` has type `bool`. The only other `Val` object is a `ProcVal`. We need a way to give a type to a `ProcVal`.

When we define a procedure, we *declare* the types of each of the procedure's formal parameters and the procedure's return type. We show how to do this shortly. Once we know these declared types, we can define the type of the procedure.

We use the notation

$$[\ t_1\ ,\ t_2\ ,\ \cdots\ ,\ t_n\ =>\ t\ ]$$

to represent the type of a procedure that has $n$ formal parameters of types $t_1, t_2, \cdots t_n$, respectively, and that returns a value of type $t$.

To simplify things, we do not permit the creation of new (named) types, so every type is either an integer, boolean, or procedure.

# Types (continued)

A *type* is one of the two primitive types or a procedure type. A *type expression* is a syntactic category in our typed language that describes a type:

```
<typeExp>:PrimTypeExp   ::= <primType>
                         PrimTypeExp(PrimType primType)
<typeExp>:ProcTypeExp   ::= LBRACK <typeExps> RARROW <typeExp> RBRACK
                         ProcTypeExp(TypeExps typeExps, TypeExp typeExp)
<primType>:IntPrimType ::= INT
                         IntPrimType()
<primType>:BoolPrimType::= BOOL
                         BoolPrimType()
<typeExps>              **= <typeExp> +COMMA
                         TypeExps(List<TypeExp> typeExpList)
```

Here are some examples of type expressions:

1. `int`

2. `bool`

3. `[ bool, int => int ]`

4. `[ [int=>bool], [int,int=>int] => [int,int=>bool] ]`

5. `[ => int ]`

Type expressions 1 and 2 describe integer and boolean types, respectively. Type expression 3 describes a procedure type that takes two parameters of type boolean and integer, respectively, and that returns an integer. Type expression 4 describes a procedure type that takes two procedure parameters of type `[int=>bool]` and `[int,int=>int]`, respectively, and that returns a procedure of type `[int,int=>bool]`. Type expression 5 describes a procedure type that takes no parameters and that returns an integer.

Our typed language now requires that every procedure expression be annotated with type expressions for each of the procedure's formal parameters – its *(declared) formal parameter types* – and a type expression for the procedure's return value – its *(declared) return type*. Together, these annotations completely define the procedure's type, called (not surprisingly) the *defined type* of the procedure.

# **Types** (continued)

The following examples show how to annotate a procedure expression in the SET language so that it conforms to the TYPE0 language rules (see Slide 9 for the start of the discussion of the TYPE0 language):

```
% SET language, no type annotations
proc(t) +(t,5)
% the same procedure in the TYPE0 language with type annotations
proc(t:int):int +(t,5)
   % defined type [int=>int], return type int


% SET language, no type annotations
proc(f,x) .f(x)
% the same procedure in the TYPE0 language with type annotations
proc(f:[int=>int], x:int):int .f(x)
   % defined type [[int=>int],int=>int], return type int
```

When used in an expression, a procedure *application* is considered to have a type equal to the *return type* of the procedure. For example, the type of the following procedure application is int:

```
.proc(t:int):int +(t,5) (4) % evaluates to 9, an int
```

We will shortly define grammar rules for the TYPE0 language that parse expressions such as these.

# **Types** (continued)

*Every expression has a type. The type of a variable (`VarExp`) is the type that is bound to the variable in the current environment.*

A *type error* is one of the following:

- an attempt to define a procedure whose declared return type does not match the type of the body of the procedure;
- an attempt to apply a non-procedure as a procedure;
- an attempt to apply a procedure to the wrong number of actual parameters;
- an attempt to apply a procedure to actual parameters whose types do not match the declared types of the corresponding formal parameters;
- an attempt to use a non-boolean as the test in an `if` statement;
- an attempt to have expressions of different types in the `then` and `else` parts of an `if` statement;
- an attempt to assign a value to a variable with a `set` statement where the type of the variable does not match the type of the RHS expression.

In the above, we consider primitive procedures such as `+`, `add1`, and `zero?` as procedures for the purposes of type checking. These primitive procedures have pre-defined procedure types.

Our static type checking implementation also detects errors where an identifier is referred to in a particular expression without appearing on the LHS of an enclosing `let` or `letrec` or in a top-level `define` – *i.e.*, if the identifier occurs free in the expresion. It also detects errors where the same identifier appears twice on the LHS of a `let` or `letrec` or if the same identifier appears twice in the formal parameter list of a `proc`.

A variable is in *declaration position* if it occurs on the LHS of a `define`, `let`, or `letrec` definition, or if it appears as a formal parameter in a procedure definition. These occurrences determine the *defined type* of the variable according to the following rules:

- The defined type of a variable appearing in declaration position in a `define`, `let`, or `letrec` is the type of its RHS expression.

- The defined type of a variable appearing in a formal parameter list is its declared type in the procedure definition.

When a variable appears in an expression, the type of the variable is its defined type using static scope rules.

# Types (continued)

Notice that static scope rules make it possible to determine the types of variables that appear free in a procedure body. With dynamic scope rules, the type of such a variable may differ from one procedure application to another, making compile-time type checking impossible.

Numeric literals all have type `int`. We predefine the identifiers `true` and `false` to have type `bool`, with the obvious semantics in the test expression of `if...` expressions. The type of a `let` or `letrec` expression is the type of the expression's body.

Since the type of a variable is determined by static scope rules, the only "interesting" types are those of procedures. We need to add syntax rules that permit us to declare the return type of a procedure and the types of its formal parameters.

# Language TYPE0

We define a language named `TYPE0` that is based on the `SET` language and that implements type syntax. We first require that a procedure definition – in the `<proc>` grammar rule – declare the types of each of its formal parameters and the return type of the procedure. Each of these declarations are given by a type expression, described earlier, in the following grammar rules related to a `<proc>`. Notice that these grammar rules are similar to the original grammar rules for the `SET` language, with the exception that they include type declarations.

```
<proc>      ::= PROC LPAREN <formals> RPAREN COLON <typeExp> <exp>
                Proc(Formals formals, TypeExp typeExp, Exp exp)
<formals>   **= <VAR> COLON <typeExp> +COMMA
                Formals(List<Token> varList, List<TypeExp> typeExpList)
```

These changes do *not* affect the behavior of the `eval` methods that define the semantics of the `SET` language. Thus the resulting `TYPE0` language has the same semantics as the `SET` language, except for the presence of boolean values: the type information in a program is simply ignored.

# Language TYPE0 (continued)

A procedure definition in our `TYPE0` language now looks like this, for example:

```
define add2 = proc(x:int):int add1(add1(x))
```

This says that the *declared type* of the formal parameter `x` is `int` and that the *declared return type* of `add2` is `int`, so that the *defined type* of `add2` is `[int=>int]`.

Since we have made changes to our syntax, our expression parse trees have additional structure, but because our `eval` methods are not affected by the type information, the `eval` methods in the `code` file for the `SET` language are entirely unchanged in the `code` file for `TYPE0`.

Notice, however, that we have not implemented any of our type rules, so no type checking is taking place.

# Language TYPE0 (continued)

Our `TYPE0` language has two new expressions: a `TrueExp` and a `FalseExp`, with obvious concrete and abstract syntax:

```
<exp>:TrueExp    ::= TRUE
                 TrueExp()
<exp>:FalseExp ::= FALSE
                 FalseExp()
```

What should the `eval` method return for these classes? So far, we have used an `IntVal` of zero to be false and everything else to be true. Once we have syntax for `true` and `false` expressions, we should introduce corresponding semantics. We therefore create a new `Val` type, `BoolVal`, that represent the values of these expressions. The code for the `BoolVal` class is in the `val` file.

Once we have a `BoolVal`, we can define the `eval` behavior for the `TrueExp` and `FalseExp` classes. Here's the code for `TrueExp` – the code for `FalseExp` is nearly identical:

```
public Val eval(Env env) {
    return new BoolVal(true);
}
```

Obviously the `isTrue` method applied to a `BoolVal` object constructed with `true` should return true, and similarly a `BoolVal` object constructed with `false` should return false. The `isTrue` method applied to any other value (either an `IntVal` or a `ProcVal`) should throw an exception.

# Language TYPE0 (continued)

We can add a number of new primitives applied to integer arguments that return `BoolVal` values, such as the following, with obvious meanings:

```
<?
<=?
>?
>=?
=?
<>?
```

The purpose of the `?` at the end of these is to remind you that they are testing something and that they return boolean values (they are *predicates*). We also change the value returned by the `zero?` primitive to a `BoolVal`.

The following expressions evaluate as indicated:

```
<?(3,3)
  % => false (actually, a false BoolVal)
<=?(3,3)
  % => true
<>?(x,y)
  % => depends on the current bindings of x and y
=?(proc(t) t, proc(u) u)
  % => exception -- can't apply to procedures
if 1 then 2 else 3
  % => exception -- 1 is not a boolean
```

# Language TYPE1

Now that we have built the syntax for type expressions and incorporated it into our grammar, we begin to implement type checking. Our strategy is to determine the type of any expression, making type checks along the way.

Our first step is to define `Type` as a *semantic* entity in our implementation. We implement `Type` as an abstract class with three subclasses:

```
IntType
BoolType
ProcType
```

These correspond to the three classes that extend the `Val` class. Observe that `Val` objects are used to evaluate expressions, whereas `Type` objects are used for type checking. They serve similar purposes but are used in different ways.

Neither `BoolType` nor `IntType` has any instance variables. A `ProcType` object has two instance variables, as described here:

```
ProcType(List<Type> paramTypeList, Type returnType)
```

An expression is of type `IntType` if it evaluates to an integer. Similarly for a `BoolType`.

When we evaluate an expression, we use an environment to determine its expressed value. When we are doing type checking, we don't care what the expressed value actually is: only its type matters.

**Language TYPE1** (continued)

Our approach is to evaluate the type (`Type`) of an expression, doing type checks of sub-expressions along the way. If there are no type errors, we can then evaluate the value (`Val`) of the expression and return this value as we have done before.

First, when we encounter a variable in an expression, we want to determine the type of that variable. If the variable is a formal parameter in the body of a procedure, its type is determined by the type declaration of the formal parameter. For any other variable, its type can be determined using static scope rules. This suggests that we have a way of binding the variable to its type during type checking in a way that is similar to binding the variable to its value during expression evaluation. We achieve this by creating a *type environment* that is almost identical to a value environment that we have been using so far, except that our type environment returns the type bound to a variable rather than its value.

The appropriate classes are shown on the next slide.

# Language TYPE1 (continued)

```
abstract class TypeEnv
    public static TypeEnv emptyTypeEnv() // returns an empty environment
    public abstract Type applyTypeEnv(String sym)
    public abstract TypeEnv extendTypeEnv(TypeBindings typeBindings)
    public abstract TypeEnv extendTypeEnv(List<String> idList, List<Type> typeList)

class TypeEnvNull extends TypeEnv  // empty type environment class

class TypeEnvNode extends TypeEnv
    public TypeBindings typeBindings // local bindings
    public TypeEnv typeEnv      // next environment

class TypeBindings
    public List<TypeBinding> typeBindingList

class TypeBinding
    public String sym
    public Type type

abstract class Type

class IntType extends Type
class BoolType extends Type
class ProcType extends Type
```

# Language TYPE1 (continued)

The definitions of the classes relating to type environments is given in the `tenv` file.

Since most of our type checking rules relate to type equality, we must define what it means for two types to be equal. There are only three basic types, and for each of them we define a `checkEquals` method. This `void` method silently returns if the types are the same, otherwise it calls the static `typeMismatch` method in the `Type` class; this method throws a runtime exception with an appropriate message.

The two primitive types are simple. For the `BoolType` class, the `checkEquals` method is defined as follows:

```
public void checkEquals(Type t) {
    t.checkBoolType(this);
}

public void checkBoolType(BoolType t) {
    // if we get here, this must be a BoolType
}
```

The default `checkBoolType` method in the `Type` class – defined for all other types except for `BoolType` – throws a type mismatch exception.

A similar definition works for the `IntType`.

We are left to define `checkEquals` for the `ProcType` class. This is easy: two `ProcType` objects are equal if they both have the same number of formal parameter types and these formal parameter types are pairwise equal (in the proper order), and if they both have the same return type.

# Language TYPE1 (continued)

Here is the definition of `checkEquals` in the `ProcType` class:

```
public void checkEquals(Type t) {
    t.checkProcType(this);
}

// check to see if the type of the ProcType object t is the same
// as this ProcType object
public void checkProcType(ProcType t) {
    // first check the return types
    this.returnType.checkEquals(t.returnType);
    // then check the types of the formal parameters
    checkEqualTypes(this.paramTypeList, t.paramTypeList);
}
```

The `checkEqualTypes` static method in the `Type` class is straight-forward: first check for equal sizes; then iterate through both of the lists, checking pairwise for type equality.

```
public static void checkEqualTypes(List<Type> t1List, List<Type> t2List) {
    if (t1List.size() != t2List.size())
        throw new RuntimeException("argument number mismatch");
    Iterator<Type> t1i = t1List.iterator();
    Iterator<Type> t2i = t2List.iterator();
    while (t1i.hasNext()) {
        Type t1 = t1i.next();
        Type t2 = t2i.next();
        t1.checkEquals(t2);
    }
}
```

# Language TYPE1 (continued)

At this point we can also add primitives operations for `and`, `or`, and `not`.

```
<prim>AndPrim ::= AND
<prim>OrPrim  ::= OR
<prim>NotPrim ::= NOT
```

Here is the `apply` code for an `AndPrim` object:

```
public Val apply(Val [] va) {
    if (va.length != 2)
        throw new RuntimeException("two arguments expected");
    boolean b0 = va[0].boolVal();
    boolean b1 = va[1].boolVal();
    return new BoolVal(b0 && b1);
}
```

It's easy to see how the `apply` code should be implemented for the other two primitives.

Although primitives are not procedures (you can't bind a variable to `add1`, for example), they do have specific type behaviors that can be described in terms of procedure types. The + primitive behaves like a `[int,int=>int]`, and the `add1` primitive behaves like a `[int=>int]`. Each of the primitives are associated with a `ProcType` that is used for type checking. We define these types in the `Type` class as static instance variables.

Rather than building each of these types by hand with constructors, we define and use a special static method `compile` in the `Type` class that takes a compact string representation of a procedure type such as `[int,int=>int]` (in this case, the string is `"ii>i"`) and returns a `ProcType` object with the proper formal parameter types and result type. You can see how the `compile` method is used in the following example definitions and their corresponding interpretations:

```
public static Type ii_i = compile("ii>i"); // [int,int=>int]
public static Type i_i  = compile("i>i");  // [int=>int]
public static Type ii_b = compile("ii>b"); // [int,int=>bool]
public static Type bb_b = compile("bb>b"); // [bool,bool=>bool]
```

The `evalType` method for the `PrimappExp` class asks the primitive object what its type is, and the object returns the appropriate type from among the ones compiled in the `Type` class. For example, in the `AddPrim` class, we define:

```
public Type toType() {
    return Type.ii_i;
}
```

# Language TYPE1 (continued)

When processing declaring the formal parameter types of procedures and process-ing `let` and `letrec` expressions, we want to ensure that there are no duplicate variable names. This amounts to checking the `varList` instance variables in the `Formals`, `LetDecls`, and `LetrecDecls` classes. This must be done after pars-ing, since the parser only uses the grammar rules and these rules do not have any way to check for duplicates. We check for duplicates during type checking and be-fore expression evaluation.

One way to check for duplicates is to use a Java `Map` object: for each variable in the list of variables, if it's already in the map, return an error, otherwise add it to the map. Here is the code:

```java
public void checkDuplicates(List<String> varList) {
    Map<String,String> map = new HashMap<String,String>();
    for (String id : varList) {
        String s = map.get(id);
        if (s != null)
            throw new RuntimeException("duplicate identifier: "+id);
        map.put(id, id);
    }
}
```

**Language TYPE1** (continued)

While the type environment classes are similar to those for expression environments, it is unnecessary to use references, since we never modify the type of a variable. Therefore type environments are similar to the environment implementation in language `V6`. The essential changes are to replace `Var` with `Type`, `Env` with `TypeEnv`, and `Binding` with `TypeBinding`.

We are now prepared to write the code to type-check expressions. We need to initialize a top-level type environment, which we call `initTypeEnv`. This happens in the `Program` class, where we also define the top-level expression environment.

```
static TypeEnv tenv = TypeEnv.initTypeEnv();
```

In the `Eval` class's `toString` method, we type-check by applying the `evalType` on the `exp` object – the "program" that we are to evaluate. The actual return value of this is irrelevant, since its purpose is simply to carry out type checking prior to expression evaluation. If the type checking succeeds, the `toString` method returns the string representation of the value of the expression.

```
public String toString() {
    exp.evalType(Program.tenv); // type check, then ...
    return exp.eval(Program.env).toString(); // ... evaluate
}
```

# Language TYPE1 (continued)

The type of a `LitExp` is easy:

```
public Type evalType(TypeEnv tenv) {
    return Type.intType; // a singleton in the Type class
}
```

So are the types of a `TrueExp` and a `FalseExp`. For both the `TrueExp` and `FalseExp` classes, we have

```
public Type evalType(TypeEnv tenv) {
    return Type.boolType;
}
```

The type of a `VarExp` looks up the type of the variable in the type environment and returns the result:

```
public Type evalType(TypeEnv tenv) {
    return tenv.applyTypeEnv(var);
}
```

**Language TYPE1** (continued)

The type of a `LetExp` involves checking the LHS variable list for duplicates, determining the types of the RHS expressions using the current type environment, extending the current type environment by binding the variables to the types of the RHS expressions, and evaluating the type of the `let` body using this extended environment.

```
public Type evalType(TypeEnv tenv) {
    TypeEnv ntenv = letDecls.addTypeBindings(tenv); // extended tenv
    return exp.evalType(ntenv);
}
```

Most of the work is done in the `addTypeBindings` method in the `LetDecls` class, which we show here.

```
public TypeEnv addTypeBindings(TypeEnv tenv) {
    List<String> idList = new ArrayList<String>(); // LHS vars
    for(Token t : varList) // LHS tokens
        idList.add(t.toString());
    TypeEnv.checkDuplicates(idList);
    Rands rands = new Rands(expList);   // RHS expressions
    List<Type> typeList = rands.evalTypeRands(tenv); // RHS types
    TypeBindings typeBindings = new TypeBindings(idList, typeList);
    return tenv.extendTypeEnv(typeBindings);
}
```

# Language TYPE1 (continued)

The `Rands` object must define the `evalTypeRands` method, which is straight-forward. Notice that this is where static scope rules come into play: in a `Let` expression, the types of the RHS expressions are evaluated in the enclosing type environment, `tenv`.

```
public List<Type> evalTypeRands(TypeEnv tenv) {
    List<Type> typeList = new ArrayList<Type>();
    for (Exp e : expList)
        typeList.add(e.evalType(tenv));
    return typeList;
}
```

## Language TYPE1 (continued)

In order to determine the defined type of a procedure, we need to give semantic meaning to the type expressions used to declare the types of a procedure's formal parameters and return type. Specifically, we need to take a type expression given by the `<typeExp>` grammar rules and evaluate its `Type`: in other words, the semantics of a type expression is its `Type`. This is parallel to taking a value expression give by `<exp>` grammar rules and evaluating its `Val`. Both of these evaluations are examples of semantics.

For each type expression, we implement a `toType` method that returns the corresponding instance of `Type`. One thing to observe is that a type expression does not involve any variables, so there is no need to keep track of a type environment.

# Language TYPE1 (continued)

For example, a type expression of the form `[int=>bool]` would evaluate to a `ProcType`, and a type expression of the form `int` would evaluate to a `IntType`.

Here's the definition of `toType` for the `PrimTypeExp` class (representing either an `int` or a `bool`) in the code file:

```
PrimTypeExp
%%%
    public Type toType() {
        return primType.toType(); // intType or boolType
    }
%%%
```

For the `BoolPrimType` class, which corresponds to grammar rules for the `true` and `false` tokens, we simply return `boolType`.

```
BoolPrimType
%%%
    public Type toType() {
        return Type.boolType; // a singleton in the Type class
    }
%%%
```

Simlar remarks apply to `IntPrimType`.

# Language TYPE1 (continued)

For a `ProcTypeExp`, we

1. determine the types (using the `toType` method) of the (formal parameter) type expressions that appear to the left of the '=>' token and insert them into a `List<Type>` object,

2. determine the the type of the (return) type expresion that appears to the right of the '=>' token, and then

3. construct the appropriate `ProcType` object.

```
ProcTypeExp
%%%
    public Type toType() {
        List<Type> paramTypeList = new ArrayList<Type>();
        for (TypeExp e : typeExps.typeExpList)
            paramTypeList.add(e.toType());
        Type returnType = typeExp.toType();
        return new ProcType(paramTypeList, returnType);
    }
%%%
```

**Language TYPE1** (continued)

We are now in a position to determine the *defined type* of a `proc` expression. By defined type, we mean the type of the procedure (a `ProcType`) as determined by the declared types of each of the formal parameters and the declared return type. For example,

```
proc(x:int):bool ...
```

has a formal parameter of declared type `int` and a declared return type of `bool`, so the defined type of this procedure is `[int=>bool]` *The defined type of a procedure is not dependent on the type environment in which the procedure expression appears, nor on the names of the formal parameters.*

However, the type of the procedure *body* is dependent on the type environment in which the procedure is evaluated, since the procedure body may have variables that are free with respect to the procedure's formal parameters but that are bound to types defined in an enclosing environment. One of our type checks is to ensure that the procedure's declared return type is the same as the type of the procedure body.

The `definedType` method in the `Proc` class evaluates the procedure's defined type, which must be a `ProcType`: it gets the declared types of the formal parameters from the `Formals` object and the declared return type from the procedure's `typeExp`. The defined type is determined exclusively by the declared types of the formal parameters and the procedure's declared return type.

```
Proc
%%%
    public ProcType definedType() {
        List<Type> declaredTypeList = formals.formalTypeList();
        Type declaredReturnType = typeExp.toType();
        return new ProcType(declaredTypeList, declaredReturnType);
    }
%%%
```

The `Formals` class defines the `formalTypeList` method:

```
Formals
%%%
    public List<Type> formalTypeList() {
        List<String> stringVarList = new ArrayList<String>();
        for(Token t : varList) {
            stringVarList.add(t.toString());
        }
        TypeEnv.checkDuplicates(stringVarList);
        List<Type> typeList = new ArrayList<Type>();
        for (TypeExp te : typeExpList)
            typeList.add(te.toType());
        return typeList;
    }
%%%
```

# Language TYPE1 (continued)

We are now ready to define `evalType` in the `Proc` class.

```
public Type evalType(TypeEnv tenv) {
    // retrieve the declared return type of the proc
    Type declaredReturnType = typeExp.toType();
    // bind the formal parameters to their declared types
    TypeBindings typeBindings = formals.declaredTypeBindings();
    // extend the tenv by the formal param type bindings
    TypeEnv ntenv = tenv.extendTypeEnv(typeBindings);
    // evaluate the type of the body using this extended
    // type environment
    Type expType = exp.evalType(ntenv);
    // and check that the declared return type matches the body type
    Type.checkEquals(declaredReturnType, expType);
    // finally build the ProcType
    List<Type> formalTypeList = formals.formalTypeList();
    return new ProcType(formalTypeList, declaredReturnType);
}
```

**Language TYPE1** (continued)

Type checking for an `if` expression involves making sure that the test expression has type `bool` and that the true part and the false part have the same types. Thus, for a `IfExp`, we have:

```
public Type evalType(TypeEnv tenv) {
    Type testType = testExp.evalType(tenv);
    testType.checkBoolType(Type.boolType);
    Type trueExpType = trueExp.evalType(tenv);
    Type falseExpType = falseExp.evalType(tenv);
    Type.checkEquals(trueExpType, falseExpType);
    return trueExpType;
}
```

What is the type of a { ... } expression? Since the purpose of this statement
is normally to perform side-effects and to return only the value of the last sub-
expression, we assume that the type of '{ ... }' is the type of the last sub-
expression. The types of the previous sub-expressions are ignored except that they
must not have type errors.

Thus the expression

```
let
  x = 1
  double = proc(t:int):int set t = *(2,t)
in
  { if zero?(x) then 0 else set x = add1(x)
  ; .double(x)
  ; proc():int x
  }
```

has type [=>int ].

# Language TYPE1 (continued)

The following method implements type checking in the `SeqExp` class:

```
public Type evalType(TypeEnv tenv) {
    Type t = exp.evalType(tenv);
    for (Exp e : seqExps.expList)
        t = e.evalType(tenv);
    return t;
}
```

In a `set` expression, we look up the type of the LHS variable and make sure that it matches the type of the RHS expression. The following method is defined in the `SetExp` class:

```
public Type evalType(TypeEnv tenv) {
    Type varType = tenv.applyTypeEnv(var);
    Type expType = exp.evalType(tenv);
    Type.checkEquals(varType, expType);
    return varType;
}
```

# Language TYPE1 (continued)

A `letrec` expression is treated similar to a `let` expression, except that the types of the RHS procedures must be evaluated in an environment that has type bindings for the LHS variables. Since the defined type of a RHS procedure can be determined solely by declared types of its parameters and return type, the types of the LHS variables can be determined independent of any environment. So we can use these LHS variables to extend the type environment and then evalute the types of the RHS procedure *bodies* using this extended environment, for the purpose of type checking with their declared return types. The `addTypeBindings` method in the `LetrecDecls` class does all the work.

```java
public TypeEnv addTypeBindings(TypeEnv tenv) {
    // first check for duplicate LHS identifiers
    List<String> stringVarList = new ArrayList<String>();
    for(Token t : varList)
        stringVarList.add(t.toString());
    TypeEnv.checkDuplicates(stringVarList);
    // next get the defined procedure types in the letrec,
    // ignoring the procedure bodies for now
    List<Type> definedTypeList = new ArrayList<Type>();
    for (Proc p : procList)
        definedTypeList.add(p.definedType()); // just the defined type
    // extend tenv with bindings of the LHS vars to the RHS defined types
    TypeBindings typeBindings =
        new TypeBindings(stringVarList, definedTypeList);
    TypeEnv ntenv = tenv.extendTypeEnv(typeBindings);
    // finally check that the proc bodies are OK in the ntenv environment
    for (Proc p : procList)
        p.evalType(ntenv);
    return ntenv;
}
```

# Language TYPE1 (continued)

Once we have type checked the RHS entries and built the extended type environ-
ment with bindings for the LHS variables, we can evaluate the type of the `letrec`
body using this extended environment.

```
LetrecExp
%%%
    public Type evalType(TypeEnv tenv) {
        TypeEnv ntenv = letrecDecls.addTypeBindings(tenv);
        return exp.evalType(ntenv);
    }
%%%
```

# Language TYPE1 (continued)

Two to go: `AppExp` and `PrimappExp`. Both of these are similar in that we need to check if the formal parameter types of the operator agree in number and type with the actual parameter (operand) types.

For an `AppExp`, we must ensure that the thing we are applying is actually a procedure – it must evaluate to a `ProcType`. We then check that the actual parameter expression types match the declared types of the formal parameters. Because a `ProcType` object already knows (and has checked the validity of) its return type, the type of an application is simply the return type of the procedure.

```
AppExp
%%%
    public Type evalType(TypeEnv tenv) {
        Type tt = exp.evalType(tenv);
        tt.checkProcType(); // make sure tt is a ProcType
        ProcType pt = (ProcType)tt;
        List<Type> argTypeList = rands.evalTypeRands(tenv);
        Type.checkEqualTypes(pt.paramTypeList, argTypeList);
        return pt.returnType;
    }
%%%
```

## Language TYPE1 (continued)

For `PrimappExp`, we do the same thing, except that we simply ask the primitive to identify its type using its `definedType` method:

```
public Type evalType(TypeEnv tenv) {
    ProcType pt = prim.definedType();
    List<Type> argTypeList = rands.evalTypeRands(tenv);
    Type.checkEqualTypes(pt.paramTypeList, argTypeList);
    return pt.returnType;
}
```

Notice that the type of a primitive is defined statically. The primitive's `definedType` method just retrieves the type, which must be a `ProcType`.

# Language TYPE1 (continued)

We are left to handle top-level `defines`. While this seems to be straight-forward, we are faced with problems of unbound variables. For example, suppose we attempt to make the follwing definitions:

```
define odd? =
  proc(x:int):bool if zero?(x) then false else .even?(sub1(x))
define even? =
  proc(x:int):bool if zero?(x) then true else .odd?(sub1(x))
```

When the procedure body of `odd?` is first encountered, the variable `even?` is unbound, so the type of the body of the procedure cannot be evaluated. Any attempt to type-check `odd?` fails with a type error.

We can remedy this situation in two ways:

- forbid top-level `defines` altogether; or

- create a mechanism to deal with unbound variables that makes top-level `defines` work.

**Language TYPE1** (continued)

We take the second approach, with some restrictions.

First, we create a new top-level operator `declare` that has a syntax similar to how we declare the type of a formal parameter in a procedure. This allows us to bind a type to an identifier without actually creating a value binding for the identifier.

Pascal provides such a mechanism with *forward* declarations, as does C/C++ with procedure declarations. Java does not suffer from this problem, in part because the order of declaration of fields and methods is not significant.

A declaration has the following concrete and abstract syntax:

```
<program>:Declare  ::= DECLARE <VAR> COLON <typeExp>
                       Declare(Token var, TypeExp typeExp)
```

# Language TYPE1 (continued)

The rules for top-level `defines` and `declares` are:

1. A top-level identifier definition serves as a declaration for the purpose of determining the type of the identifier.

2. It is an error for a top-level identifier to be declared (hence defined) more than once.

3. If an identifier declaration precedes its definition, it is an error if the declared and defined types are not identical.

4. It is an error for a declaration not to have a subsequent definition, although this error need not be checked.

Although you cannot `define` a top-level variable more than once, you can *redefine* a top-level variable using `set` if the variable has been previously `defined`. The type rule for variable assignment still requires that the RHS type of the `set` must agree with the declared type of the variable.

## Language TYPE1 (continued)

Examples of top-level declarations/definitions:

```
1.   define x = 3 % type int
     define x = 4 % error - multiple definitions for x

2.   define x = 3 % type int
     set x = 4     % ok - types agree
     set x = true % error - type of RHS must be int

3.   define fact = proc(x:int):int
        if zero?(x) then 1 else *(x,.fact(sub1(x)))
        % error - fact on the RHS has no type binding!

4.   declare fact : [int=>int]
     define fact = proc(x:int):int
        if zero?(x) then 1 else *(x,.fact(sub1(x))) %ok

5.   declare f : [=>bool]
     define y = 2
     define f = proc():int y
        % error - defined type does not match declared type

6.   declare f : [=>bool]
     define y = 2
     define f = proc():bool y
        % error - body type does not match return type
```

As a final example, consider

```
declare f : [=>int]
declare y : int
define f = proc():int y
.f() % error - no binding for y
```

The definition for f is type valid, but y has no defined value binding, so the procedure application .f() results in a run-time error.

If we subsequently give a value definition for y, the procedure application succeeds, as the following shows:

```
declare f : [=>int]
declare y : int
define f = proc():int y
define y = 3
.f() % ok - returns 3
```

We implement top-level declarations as follows:

```
Declare
%%%
    // calling toString may trigger a modificaiton
    // of the initial type environment
    public String toString() {
        TypeEnv tenv = Program.tenv; // top-level type environment
        Type tv; // variable's declared type
        String sym = var.toString(); // the LHS symbol
        try {
            // look up the variable in the initial type environment
            tv = tenv.applyTypeEnv(sym);
        } catch (RuntimeException e) {
            // no type binding -- must be a new type declaration
            // that we add to the top-level type environment
            tv = typeExp.toType();
            tenv.add(new TypeBinding(sym, tv)); // type binding
            return sym + ":" + tv;
        }
        throw new RuntimeException(sym + ": duplicate variable declaration");
    }
%%%
```

If a type binding (either through a declaration of definition) already exists, it is flagged as a duplicate declaration. Otherwise a top-level binding is made between the identifier and the type represented by its type expression. The `toString` return value for a top-level declaration/definition is the identifier being declared/defined along with its declared type.

# Language TYPE1 (continued)

We expand on top-level definitions as follows. Notice that this code is in two parts: the first part is where the variable hasn't been declared/defined before:

```
Define
%%%
    // calling toString may trigger a modification
    // of the initial type and value environments
    public String toString() {
        Env env = Program.env;        // top-level environment
        TypeEnv tenv = Program.tenv; // top-level type environment
        Type rhst; // RHS expression type
        Val  rhsv; // RHS expression value
        Type lhst; // LHS variable declared type
        Val  lhsv; // LHS variable's current value
        String sym = var.toString(); // the LHS symbol
        try {
            // look up the LHS variable in the initial type environment
            lhst = tenv.applyTypeEnv(sym);
        } catch (RuntimeException e) {
            // no type binding -- must be a new variable definition
            rhst = exp.evalType(tenv);
            tenv.add(new TypeBinding(sym, rhst));        // type binding
            rhsv = exp.eval(env);
            env.add(new Binding(sym, new ValRef(rhsv))); // value binding
            return sym + ":" + rhst; // return *something*
        }
        ...
%%%
```

# Language TYPE1 (continued)

The second part is where the variable has a declaration but possibly not a definition:

```
Define
%%%
    ...
        // the variable has a declared type, lhst -- see if it needs a def'n
        try {
            // look up the value of var in the initial environment
            lhsv = env.applyEnv(sym);
        } catch (RuntimeException e) {
            // the variable has a declared type, but no value bound to it
            // so we want to add the value binding to the top-level env.
            // first check the defined type of the RHS
            rhst = exp.evalType(tenv);
            // the declared and defined types must be the same
            Type.checkEquals(lhst, rhst);
            // get the RHS expression value
            rhsv = exp.eval(env);
            // and bind it to the variable
            env.add(new Binding(sym, new ValRef(rhsv)));
            return sym + ":" + lhst;
        }
        // the variable has a value, so must be a duplicate
        throw new RuntimeException(sym + ": duplicate variable definition");
    }
%%%
```

# Language TYPE1 (continued)

The code on the previous slide implements a top-level definition, binding a value to an identifier in the initial environment. According to our type rules, if a value binding already exists, it is flagged as an error. If a declared type binding exists without a value binding, the expression type is checked for conformance with the declared type. If a declared type binding does not exist, the top-level type environment is extended to include the type binding. In case there are no errors, the top-level value environment is extended to include the value binding.

Armed with both `declare` and `define`, we can now implement our `odd?` and `even?` procedures:

```
declare odd? : [int=>bool]
define even? =
    proc(t:int):bool if zero?(t) then true else .odd?(sub1(t))
define odd? =
    proc(t:int):bool if zero?(t) then false else .even?(sub1(t))
```